

Neural Network Model Report

Data Preprocessing

Target Variables for module include:

- APPLICATION_TYPE
- CLASSIFICATION
- IS_SUCCESSFUL

```
dtype: int64

In [5]: # Look at APPLICATION_TYPE value counts for binning
app_count = application_df["APPLICATION_TYPE"].value_counts()
app_count

Out[5]: T3      27037
        T4      1542
        T6      1216
        T5      1173
        T19     1065
        T8       737
        T7       725
        T10      528
        T9       156
        T13       66
        T12       27
        T2        16
        T25        3
        T14        3
        T29        2
        T15        2
        T17        1
        Name: APPLICATION_TYPE, dtype: int64

In [6]: # Choose a cutoff value and create a list of application types to be replaced
# use the variable name "application_types_to_replace"
application_types_to_replace = app_count[app_count<500].index

# Replace in dataframe
for app in application_types_to_replace:
    application_df["APPLICATION_TYPE"] = application_df["APPLICATION_TYPE"].replace(app,"Other")

# Check to make sure binning was successful
application_df["APPLICATION_TYPE"].value_counts()

Out[6]: T3      27037
        T4      1542
        T6      1216
        T5      1173
        T19     1065
        T8       737
        T7       725
        T10      528
        Other     276
        Name: APPLICATION_TYPE, dtype: int64
```

```
Other      276
Name: APPLICATION_TYPE, dtype: int64
```

```
In [7]: # Look at CLASSIFICATION value counts for binning
class_count = application_df["CLASSIFICATION"].value_counts()
class_count
```

```
Out[7]: C1000    17326
        C2000    6074
        C1200    4837
        C3000    1918
        C2100    1883
        ...
        C4120         1
        C8210         1
        C2561         1
        C4500         1
        C2150         1
Name: CLASSIFICATION, Length: 71, dtype: int64
```

```
In [8]: # You may find it helpful to look at CLASSIFICATION value counts >1
class_count[class_count>1]
```

```
Out[8]: C1000    17326
        C2000    6074
        C1200    4837
        C3000    1918
        C2100    1883
        C7000     777
        C1700     287
        C4000     194
        C5000     116
        C1270     114
        C2700     104
        C2800      95
        C7100      75
        C1300      58
        C1280      50
        C1230      36
        C1400      34
        C7200      32
        C2300      32
        C1240      30
        C8000      20
        C7120      18
        C1500      16
        C1800      15
        C6000      15
        C1250      14
```

```
C1267      2
C1256      2
Name: CLASSIFICATION, dtype: int64
```

```
In [9]: # Choose a cutoff value and create a list of classifications to be replaced
# use the variable name 'classifications_to_replace'
classifications_to_replace = class_count[class_count<1883].index

# Replace in dataframe
for cls in classifications_to_replace:
    application_df['CLASSIFICATION'] = application_df['CLASSIFICATION'].replace(cls,"Other")

# Check to make sure binning was successful
application_df['CLASSIFICATION'].value_counts()
```

```
Out[9]: C1000    17326
        C2000    6074
        C1200    4837
        Other    2261
        C3000    1918
        C2100    1883
Name: CLASSIFICATION, dtype: int64
```

```
In [10]: # Convert categorical data to numeric with 'pd.get_dummies'
dummies_df = pd.get_dummies(application_df)
dummies_df.head()
```

```
Out[10]:
```

	STATUS	ASK_AMT	IS_SUCCESSFUL	APPLICATION_TYPE_Other	APPLICATION_TYPE_T10	APPLICATION_TYPE_T19	APPLICATION_TYPE_T
0	1	5000	1	0	1	0	
1	1	108590	1	0	0	0	
2	1	5000	0	0	0	0	
3	1	6692	1	0	0	0	
4	1	142590	1	0	0	0	

5 rows × 44 columns

```

In [11]: # Split our preprocessed data into our features and target arrays
y=dummies_df["IS_SUCCESSFUL"].values
X=dummies_df.drop("IS_SUCCESSFUL", axis = 1)

# Split the preprocessed data into a training and testing dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state= 42)

In [12]: # Create a StandardScaler instances
scaler = StandardScaler()

# Fit the StandardScaler
X_scaler = scaler.fit(X_train)

# Scale the data
X_train_scaled = X_scaler.transform(X_train)
X_test_scaled = X_scaler.transform(X_test)

```

Features of Model include:

```

In [4]: # Determine the number of unique values in each column.
application_df.nunique()

Out[4]: APPLICATION_TYPE      17
AFFILIATION      6
CLASSIFICATION    71
USE_CASE      5
ORGANIZATION      4
STATUS      2
INCOME_AMT      9
SPECIAL_CONSIDERATIONS      2
ASK_AMT      8747
IS_SUCCESSFUL      2
dtype: int64

```

Drop Non Beneficial Information:

- EIN
- NAME

```

In [3]: # Drop the non-beneficial ID columns, 'EIN' and 'NAME'.
application_df = application_df.drop(["EIN", "NAME"], axis = 1)

```

Compiling, Training, and Evaluating the Model

How many neurons, layers, and activation functions did you select for your neural network model, and why?

- *Measuring Accuracy*

Compile, Train and Evaluate the Model

```
n [13]: # Define the model - deep neural net, i.e., the number of input features and hidden nodes for each layer.
input_layer = len(X_train_scaled[0])
hidden_layer_1= 80
hidden_layer_2 = 30

nn = tf.keras.models.Sequential()

# First hidden layer
nn.add(tf.keras.layers.Dense(units=hidden_layer_1, activation="relu", input_dim = input_layer))

# Second hidden layer
nn.add(tf.keras.layers.Dense(units=hidden_layer_2, activation="relu"))

# Output layer
nn.add(tf.keras.layers.Dense(units=1, activation="sigmoid"))

# Check the structure of the model
nn.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 80)	3520
dense_1 (Dense)	(None, 30)	2430
dense_2 (Dense)	(None, 1)	31

=====
Total params: 5,981
Trainable params: 5,981
Non-trainable params: 0

```
n [14]: # Compile the model
nn.compile(loss="binary_crossentropy", optimizer= "adam", metrics=["accuracy"])
```

Were you able to achieve the target model performance?

- No, we reached at a 72.71% accuracy

Trainable params: 5,981
Non-trainable params: 0

```
In [14]: # Compile the model
nn.compile(loss="binary_crossentropy", optimizer="adam", metrics=["accuracy"])
```

```
In [15]: # Train the model
fit_model = nn.fit(X_train_scaled, y_train, epochs=100)
```

Epoch 1/100
333/333 - 1s - loss: 0.5334 - accuracy: 0.7409

719/719 [=====] - 1s 2ms/step - loss: 0.5334 - accuracy: 0.7409

```
In [16]: # Evaluate the model using the test data
model_loss, model_accuracy = nn.evaluate(X_test_scaled, y_test, verbose=2)
print(f"Loss: {model_loss}, Accuracy: {model_accuracy}")
```

354/354 - 1s - loss: 0.5608 - accuracy: 0.7272 - 583ms/epoch - 2ms/step
Loss: 0.5607906579971313, Accuracy: 0.7271843552589417

```
In [17]: # Export our model to HDF5 file
nn.save('AlphabetSoupCharity.h5')
```

What steps did you take in your attempts to increase model performance?

Optimization Attempt#1

- Dropping more or fewer columns (dropping an additional column: organization).
- Increasing or decreasing the number of values for each bin(changed thresholds for application type in the classification column).

```
In [2]: # Drop the non-beneficial ID columns, 'EIN' and 'NAME'.
application_df = application_df.drop(["EIN", "NAME", "ORGANIZATION"], axis = 1)
```

```
In [3]: # Determine the number of unique values in each column.
application_df.nunique()
```

```
Out[3]: APPLICATION_TYPE      17
AFFILIATION                   6
CLASSIFICATION                71
USE_CASE                      5
STATUS                        2
INCOME_AMT                    9
SPECIAL_CONSIDERATIONS        2
ASK_AMT                      8747
IS_SUCCESSFUL                  2
dtype: int64
```

Name: APPLICATION_TYPE, dtype: int64

```
In [5]: # Choose a cutoff value and create a list of application types to be replaced
# use the variable `application_types_to_replace`
application_types_to_replace = app_count[app_count<50].index

# Replace in dataframe
for app in application_types_to_replace:
    application_df['APPLICATION_TYPE'] = application_df['APPLICATION_TYPE'].replace(app, "Other")

# Check to make sure binning was successful
application_df['APPLICATION_TYPE'].value_counts()
```

```
Out[5]: T3      27037
T4       1542
T6       1216
T5       1173
T19      1065
T8        737
T7        725
T10       528
T9        156
T13        66
Other       54
Name: APPLICATION_TYPE, dtype: int64
```

Optimization Attempt#2

- Add more neurons to a hidden layer(adding more neurons to the 1st and 2nd hidden layers).
- Add more hidden layers(adding a 3rd hidden layer).

```
In [16]: # Define the model - deep neural net, i.e., the number of input features and hidden nodes for each layer.
input_layer = len(X_train_scaled[0])
hidden_layer_1= 100
hidden_layer_2 = 50
hidden_layer_3 = 20

nn2 = tf.keras.models.Sequential()

# First hidden layer
nn2.add(tf.keras.layers.Dense(units=hidden_layer_1, activation="relu", input_dim = input_layer))

# Second hidden layer
nn2.add(tf.keras.layers.Dense(units=hidden_layer_2, activation="relu"))

# Third hidden layer
nn2.add(tf.keras.layers.Dense(units=hidden_layer_3, activation="relu"))

# Output layer
nn2.add(tf.keras.layers.Dense(units=1, activation="sigmoid"))

# Check the structure of the model
nn2.summary()

# Compile the model
nn2.compile(loss="binary_crossentropy", optimizer= "adam", metrics=["accuracy"])

fit_model2 = nn2.fit (X_train_scaled, y_train, epochs= 100)

model_loss, model_accuracy = nn2.evaluate(X_test_scaled,y_test,verbose=2)
print(f"Loss: {model_loss}, Accuracy: {model_accuracy}")
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
--------------	--------------	---------

Optimization Attempt #3

- Use different activation functions for the hidden layers(using sigmoid as activation function for all layers).
- Add or reduce the number of epochs to the training regimen(increasing epochs from 100 to 200).

```
In [17]: # Define the model - deep neural net, i.e., the number of input features and hidden nodes for each Layer.
input_layer = len(X_train_scaled[0])
hidden_layer_1 = 100
hidden_layer_2 = 50
hidden_layer_3 = 20

nn3 = tf.keras.models.Sequential()

# First hidden Layer
nn3.add(tf.keras.layers.Dense(units=hidden_layer_1, activation="sigmoid", input_dim = input_layer))

# Second hidden Layer
nn3.add(tf.keras.layers.Dense(units=hidden_layer_2, activation="sigmoid"))

# Third hidden Layer
nn3.add(tf.keras.layers.Dense(units=hidden_layer_3, activation="sigmoid"))

# Output Layer
nn3.add(tf.keras.layers.Dense(units=1, activation="sigmoid"))

# Check the structure of the model
nn3.summary()

# Compile the model
nn3.compile(loss="binary_crossentropy", optimizer= "adam", metrics=["accuracy"])

fit_model3 = nn3.fit (X_train_scaled, y_train, epochs= 200)

model_loss, model_accuracy = nn3.evaluate(X_test_scaled,y_test,verbose=2)
print(f"Loss: {model_loss}, Accuracy: {model_accuracy}")
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_7 (Dense)	(None, 100)	5800
dense_8 (Dense)	(None, 50)	5050
dense_9 (Dense)	(None, 20)	1020
dense_10 (Dense)	(None, 1)	21

=====

Total params: 11,891
Trainable params: 11,891
Non-trainable params: 0

Summary: My second model with adding more neurons to a hidden layer(adding more neurons to the 1st and 2nd hidden layers),and adding more hidden layers(adding a 3rd hidden layer) had the best predictive accuracy score with a 72.89.